

Explainability:

1. Model Architecture:

- The model uses a Multinomial Naive Bayes classifier, a commonly used algorithm for text classification tasks. Naive Bayes models are known for their simplicity and efficiency, making them suitable for various applications.

2. Decisions and Preprocessing:

- The dataset was loaded from an Excel file containing labeled sentences. These sentences were then shuffled to ensure a random order, which is crucial for effective model training.
- The text data was split into training, validation, and test sets, following a standard practice. The split was stratified to maintain the distribution of classes in each subset, ensuring a representative sample for training and evaluation.

3. Feature Representation:

- The model uses the TF-IDF (Term Frequency-Inverse Document Frequency) vectorization technique to represent the sentences as numerical features. TF-IDF assigns weights to words based on their frequency in a document relative to their frequency in the entire corpus.

4. Model Training:

- The Multinomial Naive Bayes classifier was trained on the TF-IDF transformed training data. Naive Bayes classifiers are well-suited for text classification tasks and are known for their speed and efficiency.

5. Evaluation on Validation Set:

- The model's performance was evaluated on a separate validation set using standard classification metrics.

Metrics:

1. Accuracy:

- Accuracy measures the overall correctness of the model. In this case, it represents the proportion of correctly classified instances out of the total instances.

2. Precision:

- Precision is the ratio of true positive predictions to the total predicted positives. It indicates how many of the predicted positive instances are actually positive. In the context of the application, precision is particularly relevant for understanding how well the model performs when it predicts the 'active' class.

3. Recall:

- Recall is the ratio of true positive predictions to the total actual positives. It provides insights into how well the model captures all instances of the positive class. In this case, it helps gauge how well the model identifies 'active' sentences.

4. F1 Score:

- The F1 score is the harmonic mean of precision and recall. It provides a balanced measure of a model's performance, especially when there is an imbalance between classes.

Analysis of Strengths and Weaknesses:

Strengths:

- The use of TF-IDF for text representation captures the importance of words in distinguishing between different classes.
- The model's simplicity and efficiency make it suitable for quick deployment and inference.

Weaknesses:

- The model's performance heavily depends on the quality and representativeness of the labeled data. If the dataset is not diverse or representative, the model might struggle with generalization.
- Naive Bayes models assume independence between features, which may not hold true for all types of data.

Recommendations for Improvement:

- Collect more diverse and representative data to enhance the model's generalization capabilities.
- Experiment with more complex models, such as deep learning architectures, to capture intricate patterns in the data.
- Consider incorporating additional features or using more advanced techniques for text representation.

Model Architecture:

3.1 Text Representation:

The model uses TF-IDF vectorization to convert sentences into numerical features. TF-IDF is chosen for its ability to capture word importance.

3.2 Model Selection:

A Multinomial Naive Bayes classifier is employed due to its simplicity and efficiency in text classification tasks. Naive Bayes is briefly explained for context.

4. Preprocessing Steps:

4.1 Data Loading:

The dataset is loaded from an Excel file containing labeled sentences.

4.2 Data Cleaning:

Unnecessary columns or rows are removed, and missing data is handled appropriately.

4.3 Data Shuffling and Splitting:

The dataset is shuffled to ensure a random order. It is then split into training (60%), validation (20%), and test (20%) sets using a stratified approach.

5. Model Evaluation:

5.1 Metrics:

The model is evaluated using metrics such as accuracy, precision, recall, and F1 score. The interpretation of each metric is explained.

5.2 Test Set Performance:

The model's performance on the test set is presented to demonstrate its generalization capabilities.

Model's Strengths:

1. Simplicity and Efficiency:

- The use of a Multinomial Naive Bayes classifier makes the model computationally efficient and easy to interpret.
- Naive Bayes models are known for their simplicity and effectiveness in text classification tasks, making them suitable for quick deployment.

2. TF-IDF Vectorization:

- The choice of TF-IDF for text representation captures the importance of words in distinguishing between active and passive voice sentences.
- TF-IDF provides a numerical representation that accounts for both the frequency of terms in a sentence and their importance across the entire dataset.

3. Stratified Data Splitting:

- The dataset is split into training, validation, and test sets using a 60:20:20 ratio with a stratified approach.

- Stratification ensures that each subset maintains the same class distribution as the original dataset, enhancing the model's ability to generalize to new data.

4. **Clear Evaluation Metrics:**

- The use of standard metrics such as accuracy, precision, recall, and F1 score provides a comprehensive evaluation of the model's performance.
- These metrics offer a balanced view of the model's ability to correctly classify both active and passive voice sentences.

Areas for Improvement:

1. **Data Diversity:**

- The model's performance heavily depends on the quality and diversity of the labeled dataset.
- Collecting a more diverse set of sentences, spanning various topics and writing styles, could enhance the model's generalization capabilities.

2. **Model Complexity:**

- While Naive Bayes models are efficient, exploring more complex models, such as deep learning architectures, might capture intricate patterns in the data.
- Experimenting with alternative algorithms could lead to improvements, especially if the relationships between words are more nuanced.

3. **Handling Imbalanced Classes:**

- If there is a significant imbalance between active and passive voice sentences, the model may benefit from techniques to handle imbalanced classes.
- Consider techniques like oversampling the minority class or using class weights during training to ensure the model doesn't become biased toward the majority class.

4. **Fine-tuning Hyperparameters:**

- Fine-tuning hyperparameters, such as the smoothing parameter in the Multinomial Naive Bayes model, might optimize the model's performance.
- A grid search or random search could be conducted to identify the optimal set of hyperparameters.

5. **Model Interpretability:**

- Enhancing model interpretability, especially for non-technical stakeholders, can be valuable. Consider using model-agnostic interpretability techniques like LIME or SHAP to explain individual predictions.

6. **Data Preprocessing Strategies:**

- Experiment with different data preprocessing strategies, such as stemming or lemmatization, to determine their impact on the model's performance.

7. Continuous Monitoring and Updating:

- Establish a system for continuous monitoring of the model's performance, and be ready to update the model as new data becomes available or the distribution of language evolves.